

METIS

INTRODUCTION TO HADOOP



What is Hadoop?



- A system for processing big data
- A framework for storing and organizing large scale data

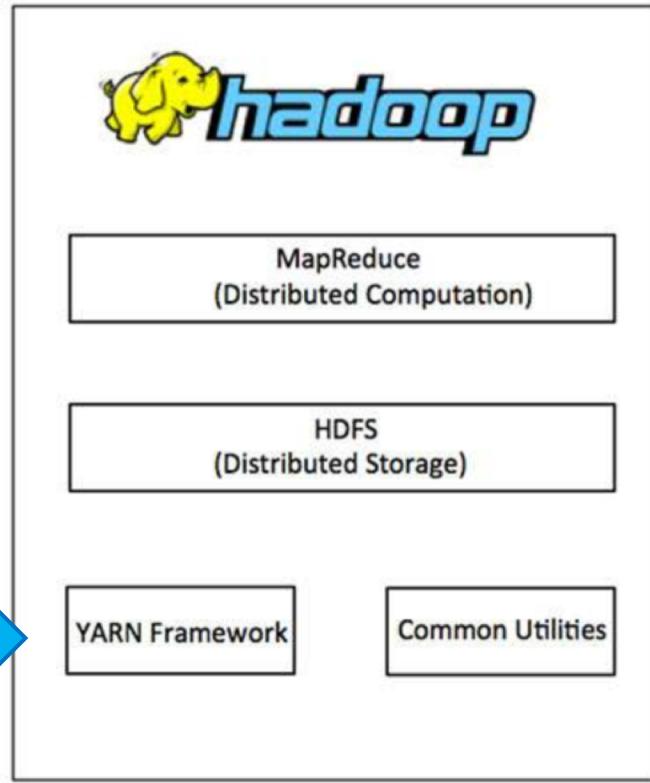
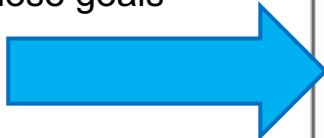


What is Hadoop?

- MapReduce
- HDFS: The Hadoop Distributed File System



Plus some other
managerial tools for
making those goals
happen



HDFS



Hadoop Distributed File System



- The largest legacy from Hadoop is HDFS
- It's a clever system for “chunking” our data up and storing it on lots of different computers for later use.
- The storage points are called “nodes,” and they are distributed across a “cluster” of computers.



We have data. When we want to use it, we just load it up and go.

```
010101010100000111101010010101010101010010101010  
1010000111111010101010110010101010101010010101  
0101000001111010100101010101010100101010101000  
0111111010101010110010101010101010101010101010  
0000111101010010101010101001010101010100001111  
110101010101100101010101010101010101010101000001  
111010100101010101010010101010101010000111111010  
101010101100101010101010101010101010100000111101  
010010101010101001010101010101000011111101010101  
01011001010101010101010101010100000111101010001  
01010101010010101010100001111110101010101011
```



So far, we've always been in this situation.

RAM
010101010100000111101010010101010101010010101010
10100001111110101010101100101010101010010101
0101000001111010100101010101010010101010101000
011111101010101011001010101010101010101010101010
0000111101010010101010101001010101010100001111
110101010101100101010101010101010101010101000001
1110101001010101010100101010101010000111111010
101010101100101010101010101010101010100000111101
0100101010101010010101010101000011111101010101
01011001010101010101010101010100000111101010001
010101010100101010101000011111101010101010111



What happens when we can't fit everything in RAM?

RAM	
0101010101000001111010100101010101010010101010	
10100001111110101010101100101010101010010101	
0101000001111010100101010101010010101010101000	
01111110101010101010010101010101010101010101010	
00001111010100101010101001010101010100001111	
110101010101011001010101010101010101010101000001	
111010100101010101010010101010101010000111111010	
101010101100101010101010101010101010100000111101	
0100101010101010010101010101000011111101010101	
010110010101010101010101010100000111101010001	
01010101010010101010100001111110101010101011	



We need a system that can handle all of this.

HDFS1

0101010101000001111
1010000111111010101
0101000001111010100
0111111010101010101

HDFS2

1111010100101010101
0101010101100101010
0100101010101010100
0101100101010101010

HDFS3

1110101001010101010
1010101011001010101
0100101010101010100
0101100101010101010

HDFS4

0101010100000111101
1000011111101010101
1010000011110101001
1111110101010101011



HDFS



- HDFS chunks data into smaller blocks (128 MB default), and distributes those blocks across a cluster
- It's designed to work on “commodity” (read: old, unreliable) hardware. So it automatically creates backups of the data across the cluster.
- It does this by using two types of “nodes” in the cluster



Name
Node

Data Node



Name Node

Data Node

- Controls where the data lives.
- Processes requests to add and find data on the cluster.
- Stores metadata like permissions, disk space, data access information.



Name Node

- Controls where the data lives.
- Processes requests to add and find data on the cluster.
- Stores metadata like permissions, disk space, data access information.

Data Node

- Stores data.
- Listens to the NameNode for commands and requests
- Stores block-data information (like verification pieces)

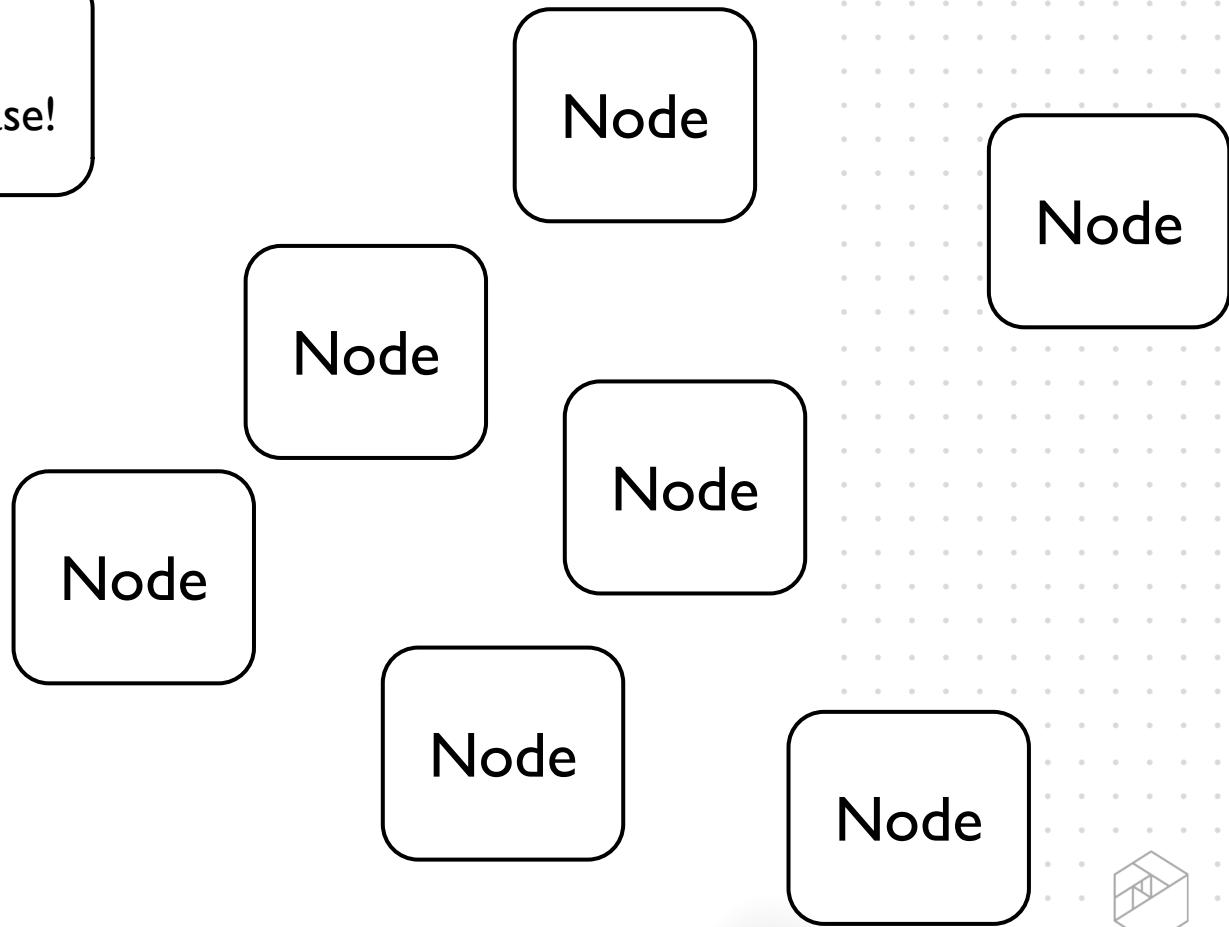


**LET'S 'BUILD' A
CLUSTER TO SEE IT ALL
IN ACTION**





I'd like a
cluster please!





I'd like a
cluster please!

Name
Node

ON IT, BOSS.

Node

Node

Node

Node

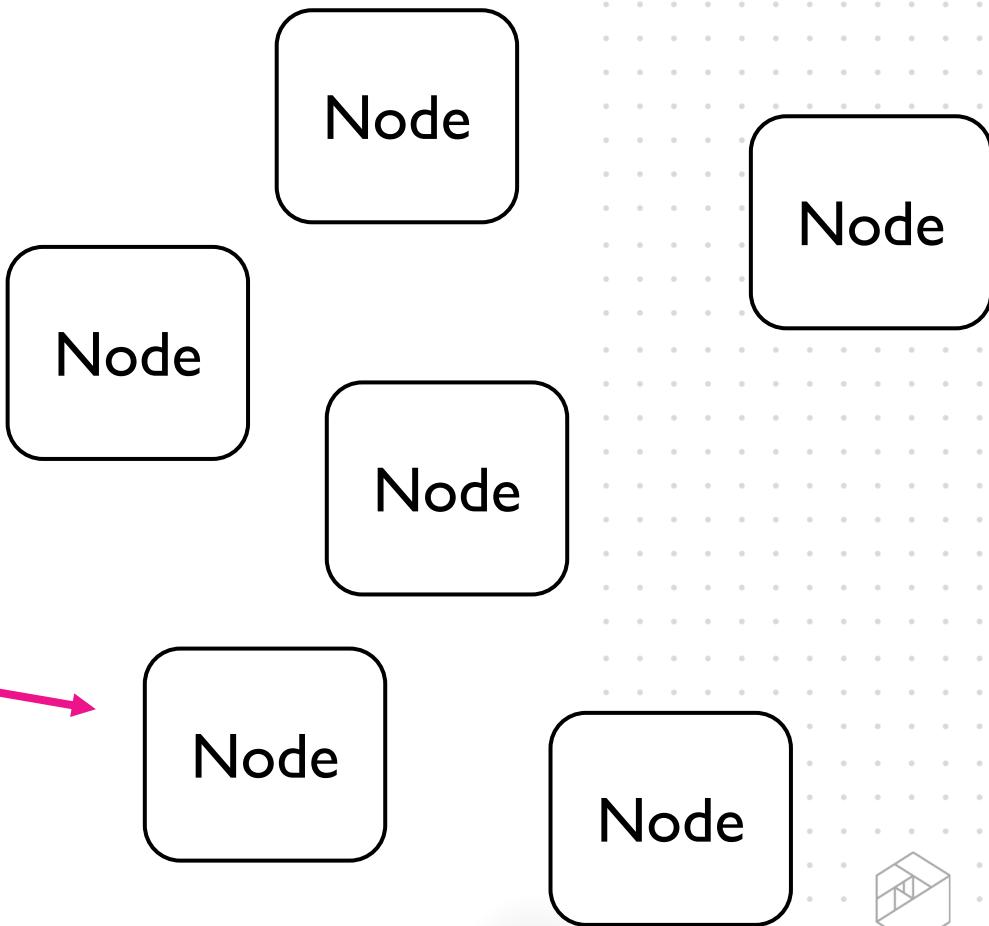
Node





HAS ANYONE CLAIMED YOU?

Name
Node





Name
Node

NOPE



Node

Node

Node

Node

Node





COOL. YOU'RE A DATANODE.

Name
Node

Data
Node

Node

Node

Node

Node

Node





HAS ANYONE CLAIMED YOU?

Name
Node

Data
Node

Node

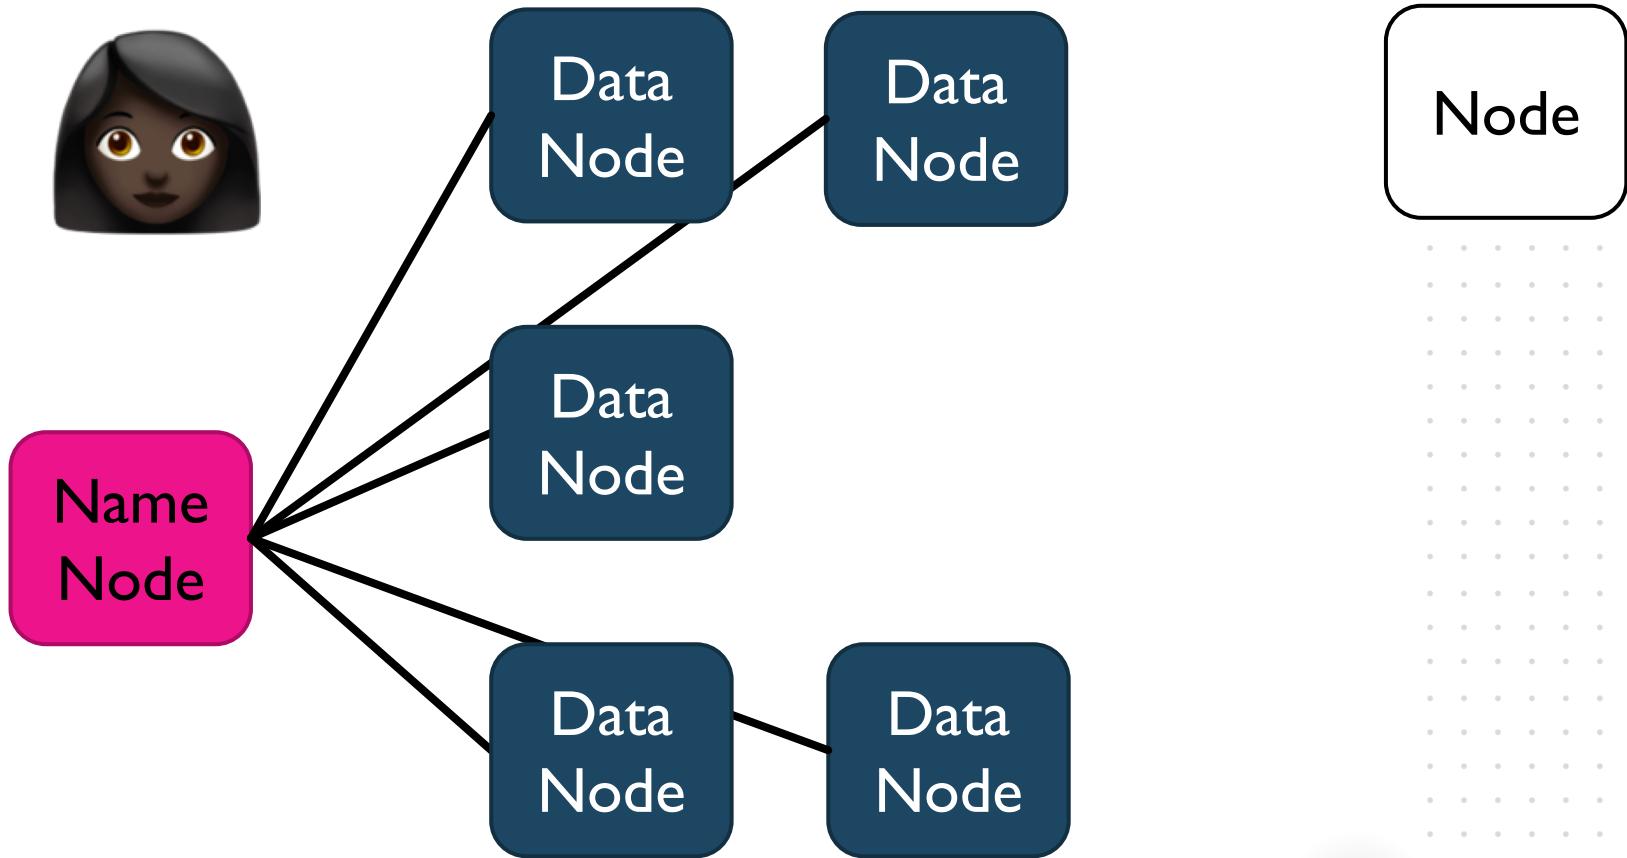
Node

Node

Node

Node



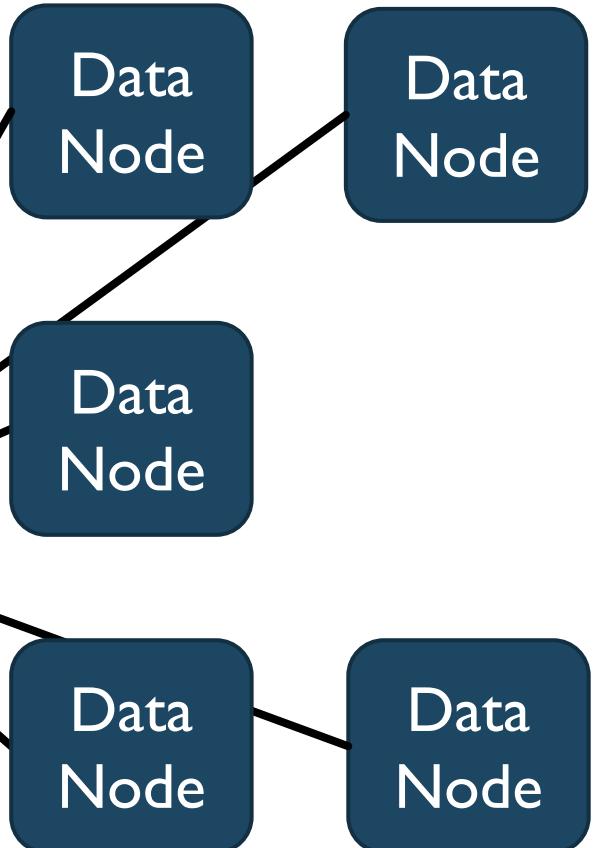


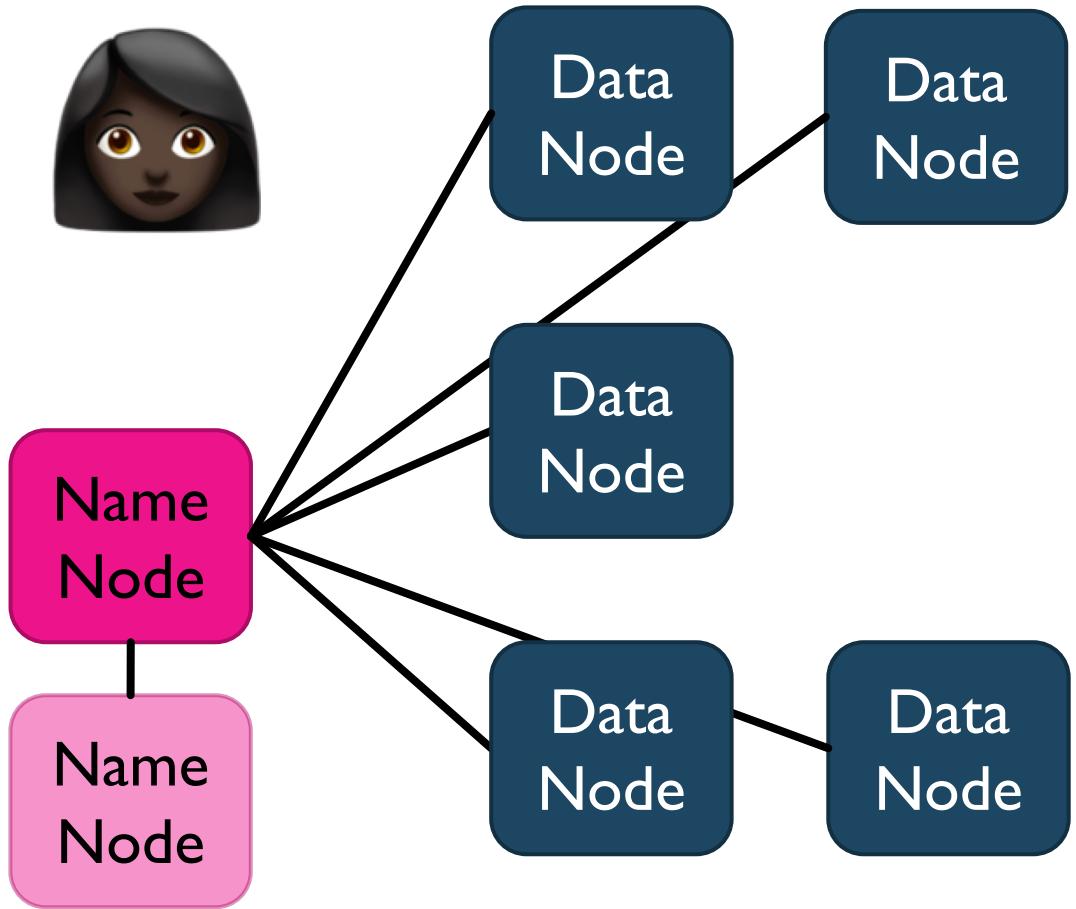


BETTER BACK ME UP.

BETTER BACK ME UP.

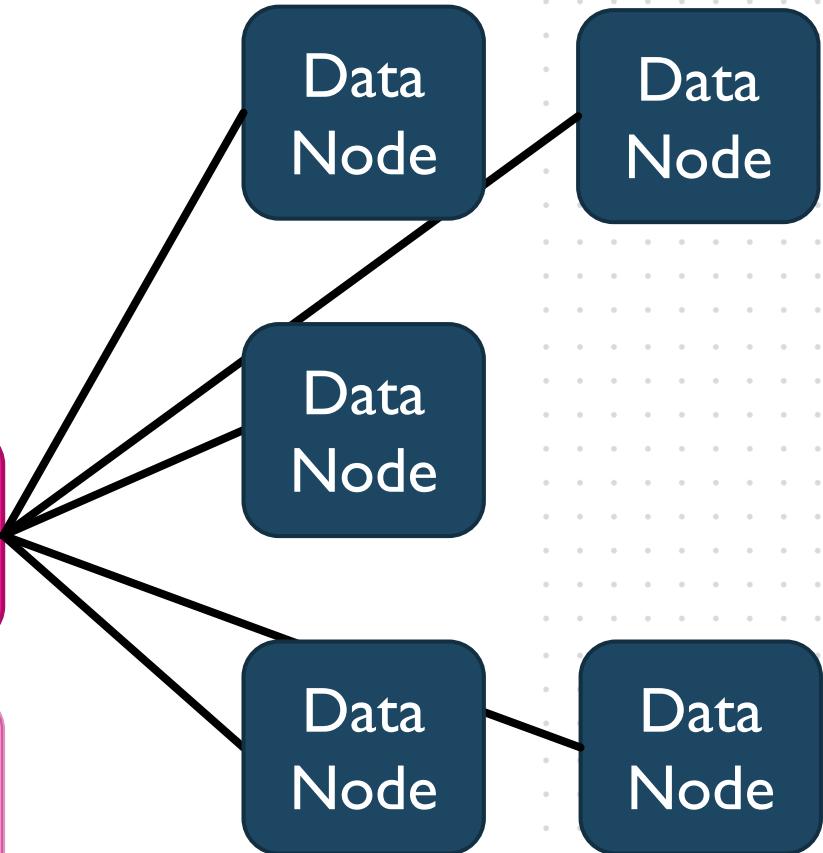
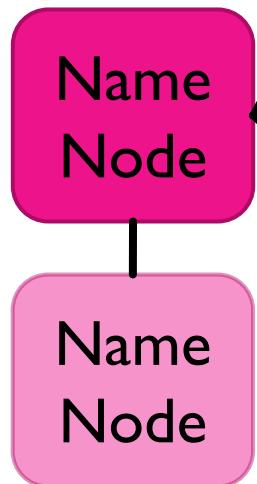
Name
Node





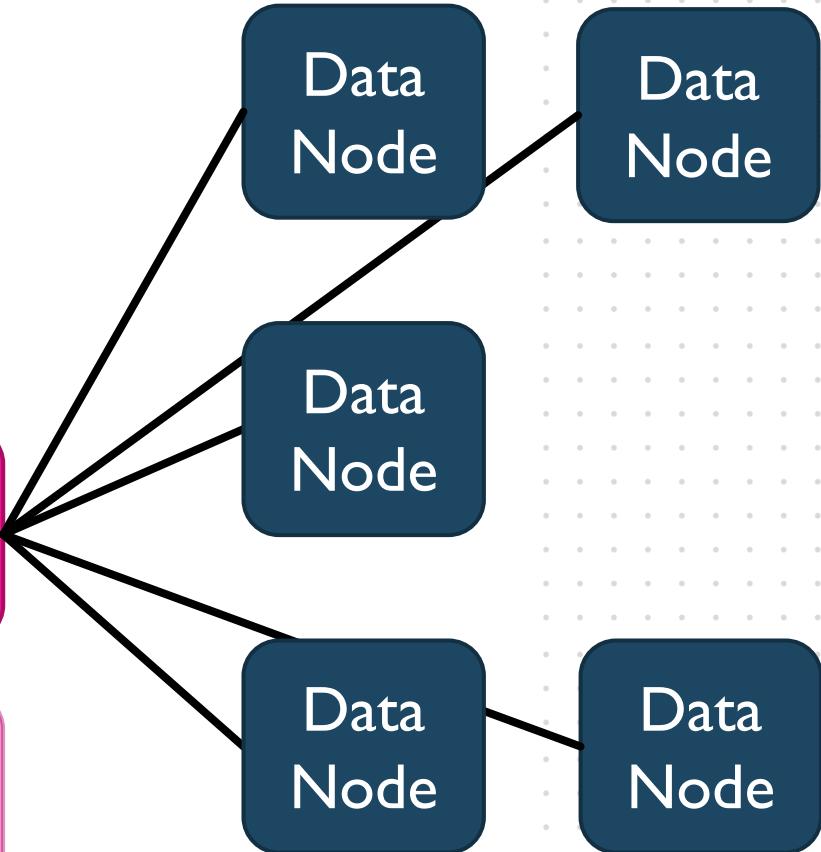
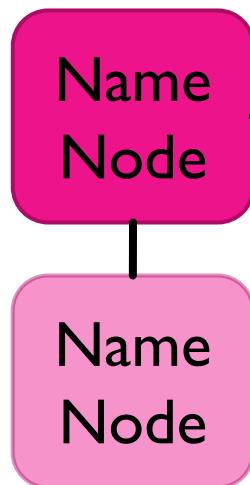


READY, BOSS.





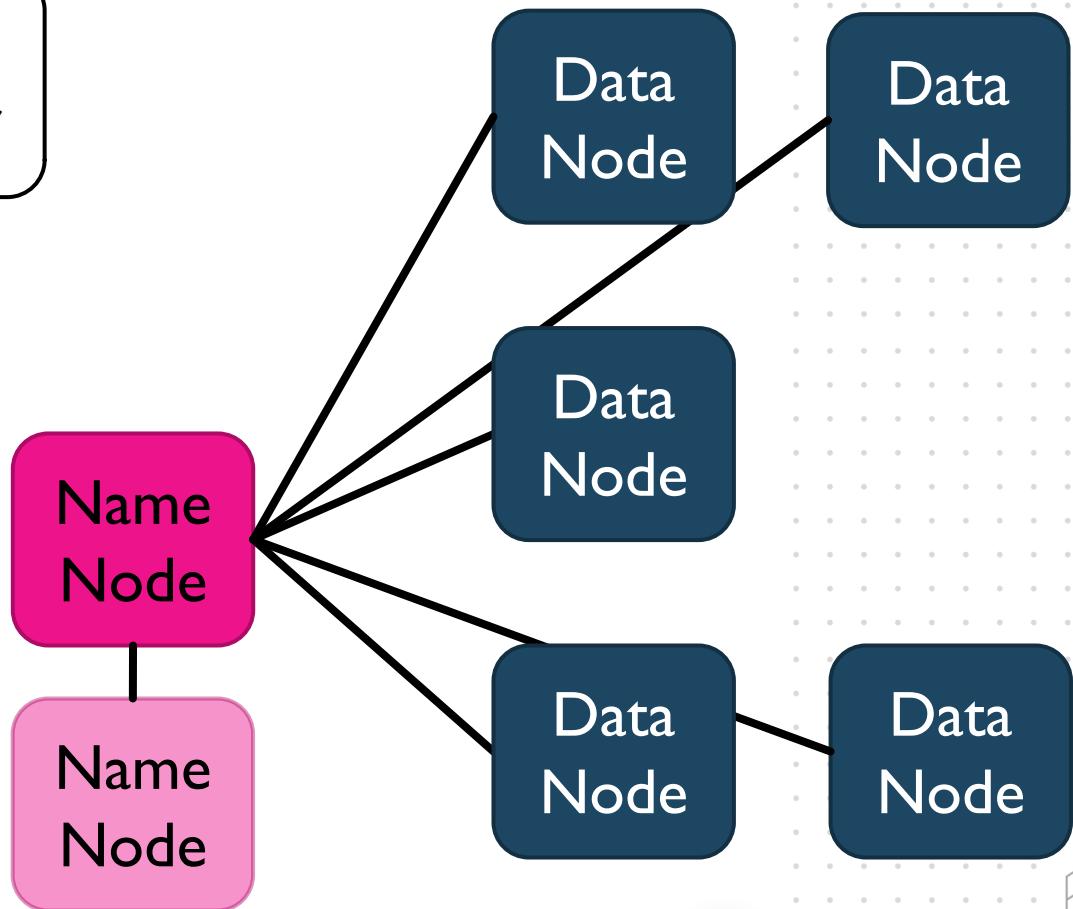
Add this data
to the cluster





Add this data
to the cluster

The system chunks the data, moves those chunks to data nodes, and then makes sure the data is backed up across many nodes.





TOO BIG. I'LL BREAK IT UP.

Name
Node

Name
Node

Data
Node

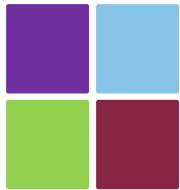
Data
Node

Data
Node

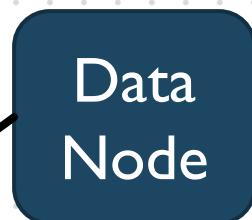
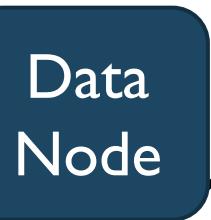
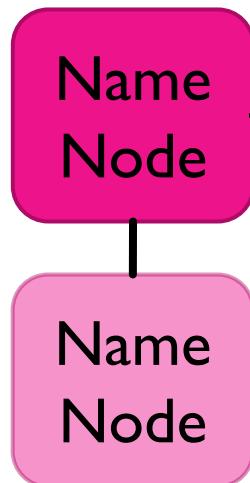
Data
Node

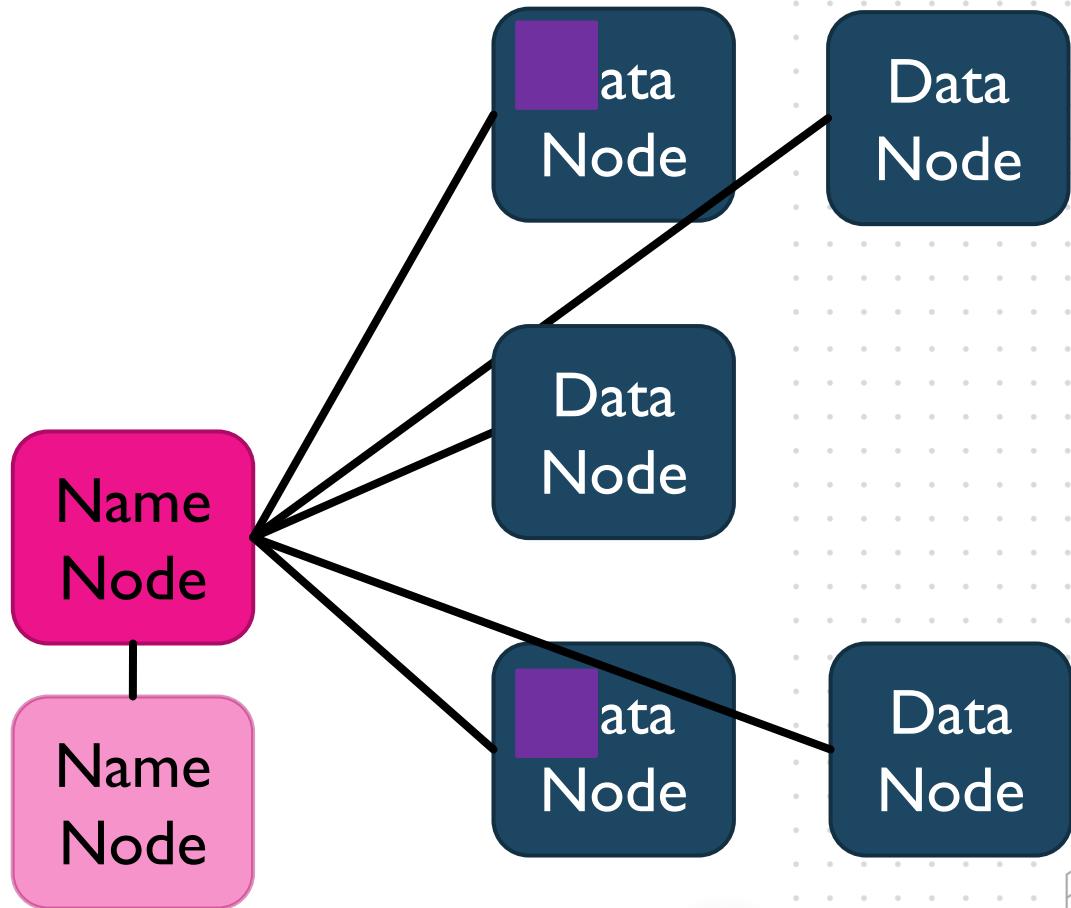
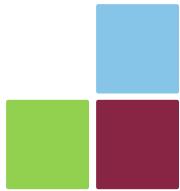
Data
Node

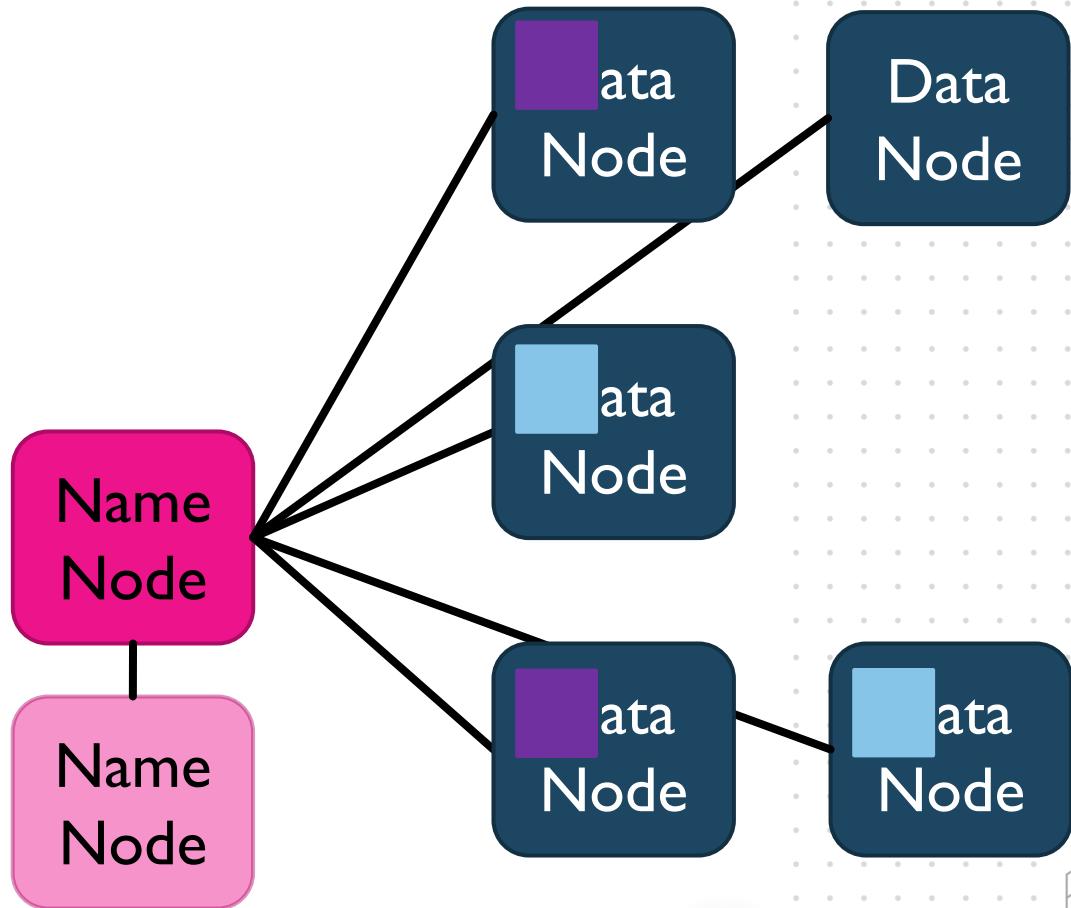


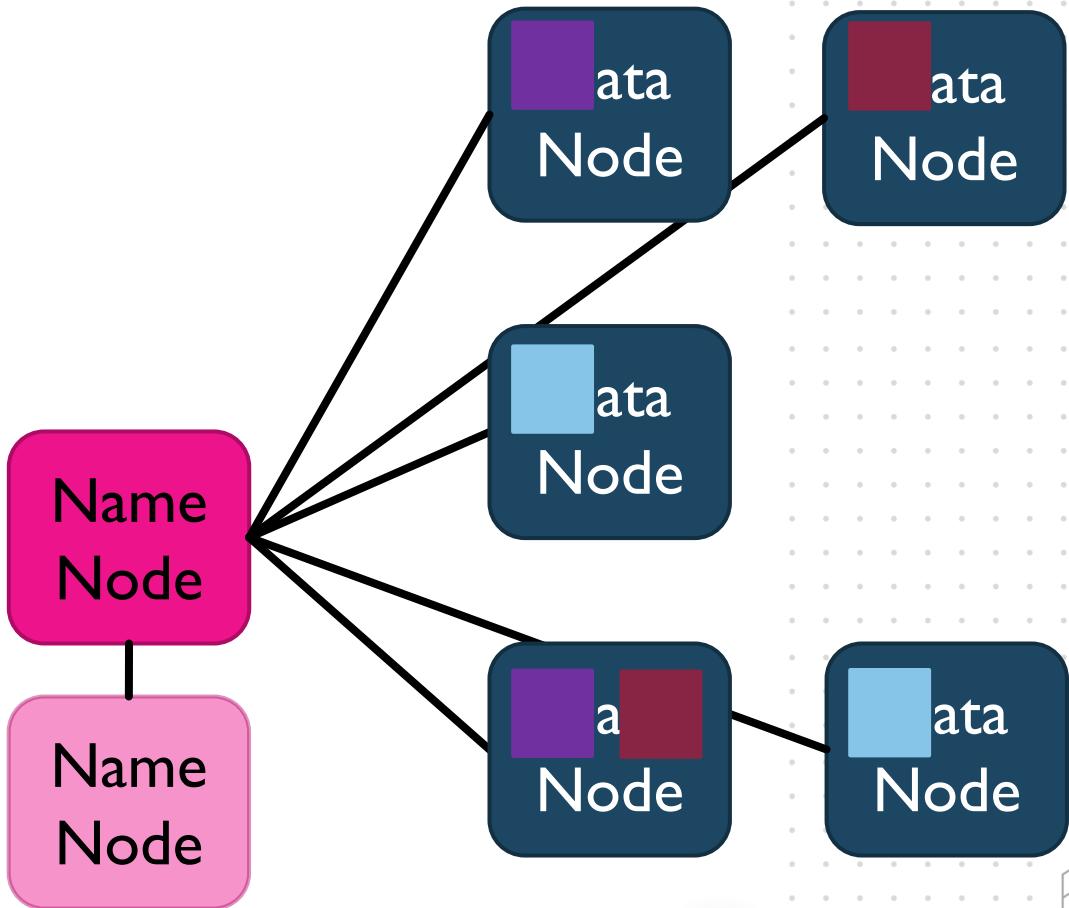


TOO BIG. I'LL BREAK IT UP.











DONE.

Name
Node

Name
Node

a
Node

a
Node

a
Node

ata
Node

ata
Node



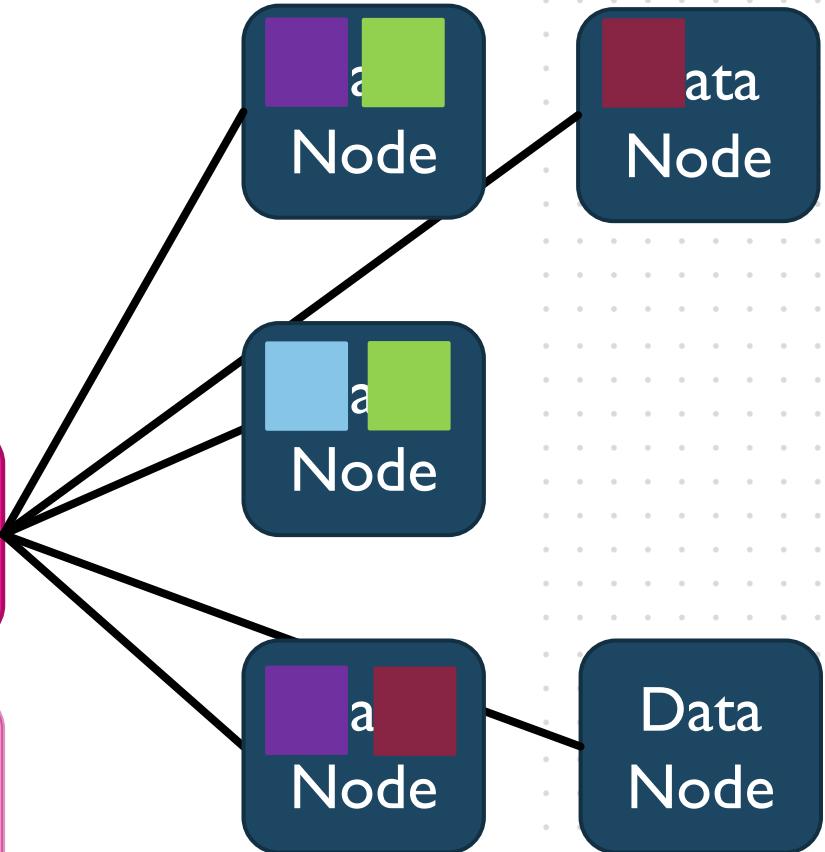
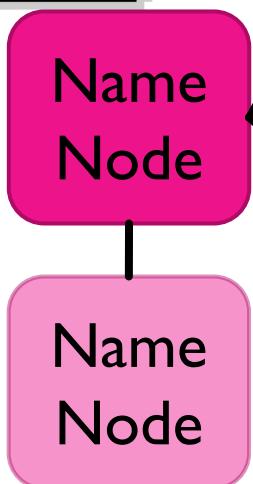
WHAT IF WE WANT TO USE THE DATA?





I'd like section
I of the data
please.

LET ME FIND IT FOR YOU.





10 ms to server

20 ms to server

Name
Node

Name
Node

a
Data
Node

a
Data
Node

a
Data
Node

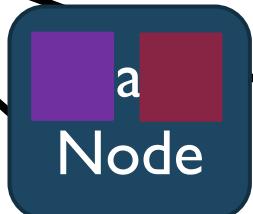
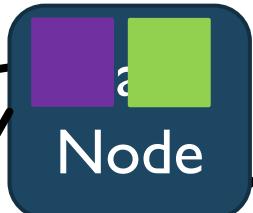
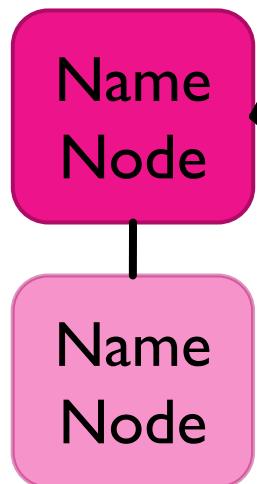
a
Data
Node

Data
Node





HERE YOU GO.



HDFS



- Redundancy
- Central control for a user
- Data management through chunking



Commodity Hardware



- It's cheaper to scale horizontally. Can always tie in more nodes from older hard drives, hard to get a 2,000,000 TB single drive.
- Allows us to salvage old systems and make a new, better system.
- Redundancy makes us fault tolerant when old hardware dies.





10 ms to server

20 ms to server

Name
Node

Name
Node

a
Data
Node

a
Data
Node

a
Data
Node

a
Data
Node

Data
Node





UHH. JUST A SECOND.

Name
Node

Name
Node

a
Node

a
Node

a
Node

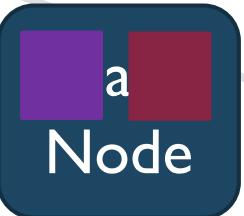
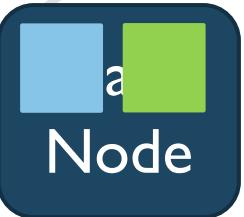
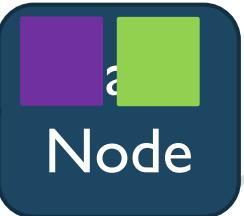
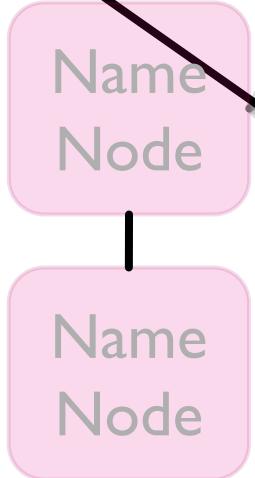
ata
Node

Data
Node





HERE YOU GO.



HDFS I/O



- HDFS supports:
 - Read
 - Write
 - Delete
- Once the data is on HDFS, it's read-only, except for appending. Makes it safe from idiots using the system.



MAPREDUCE



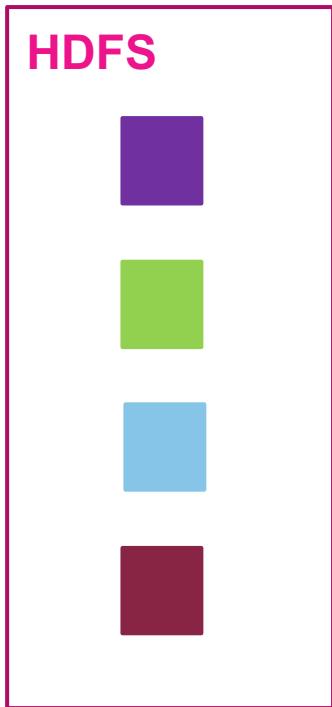
Processing Data



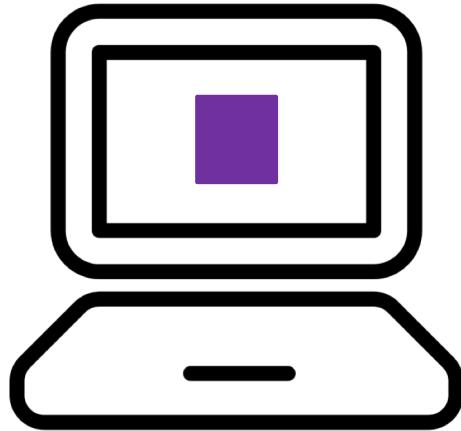
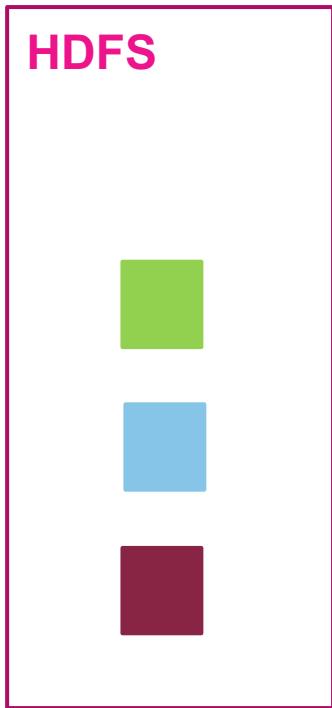
- We have two options:
 1. Pull the data from HDFS and process
 2. Make use of the cluster to process



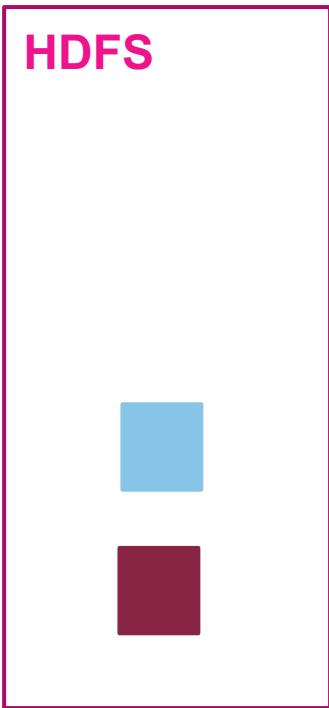
Pull-and-Process



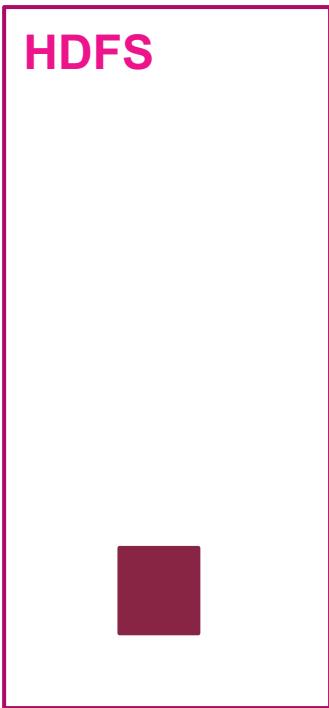
Pull-and-Process



Pull-and-Process

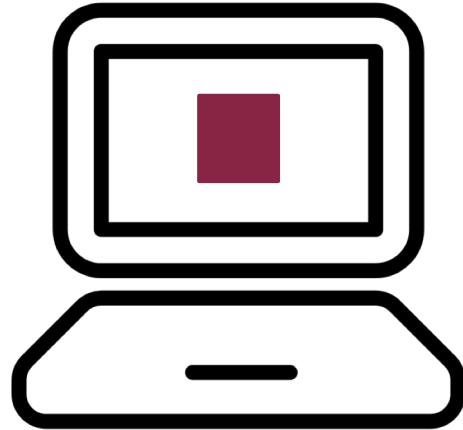


Pull-and-Process

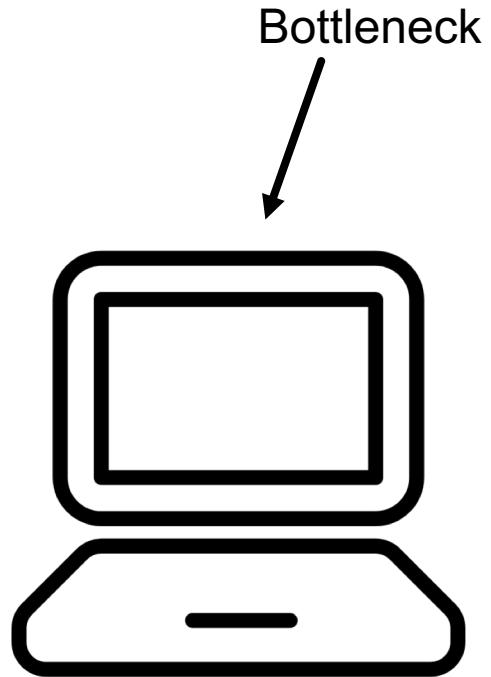
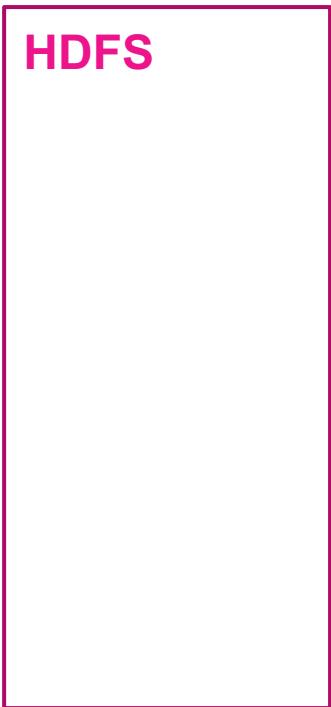


Pull-and-Process

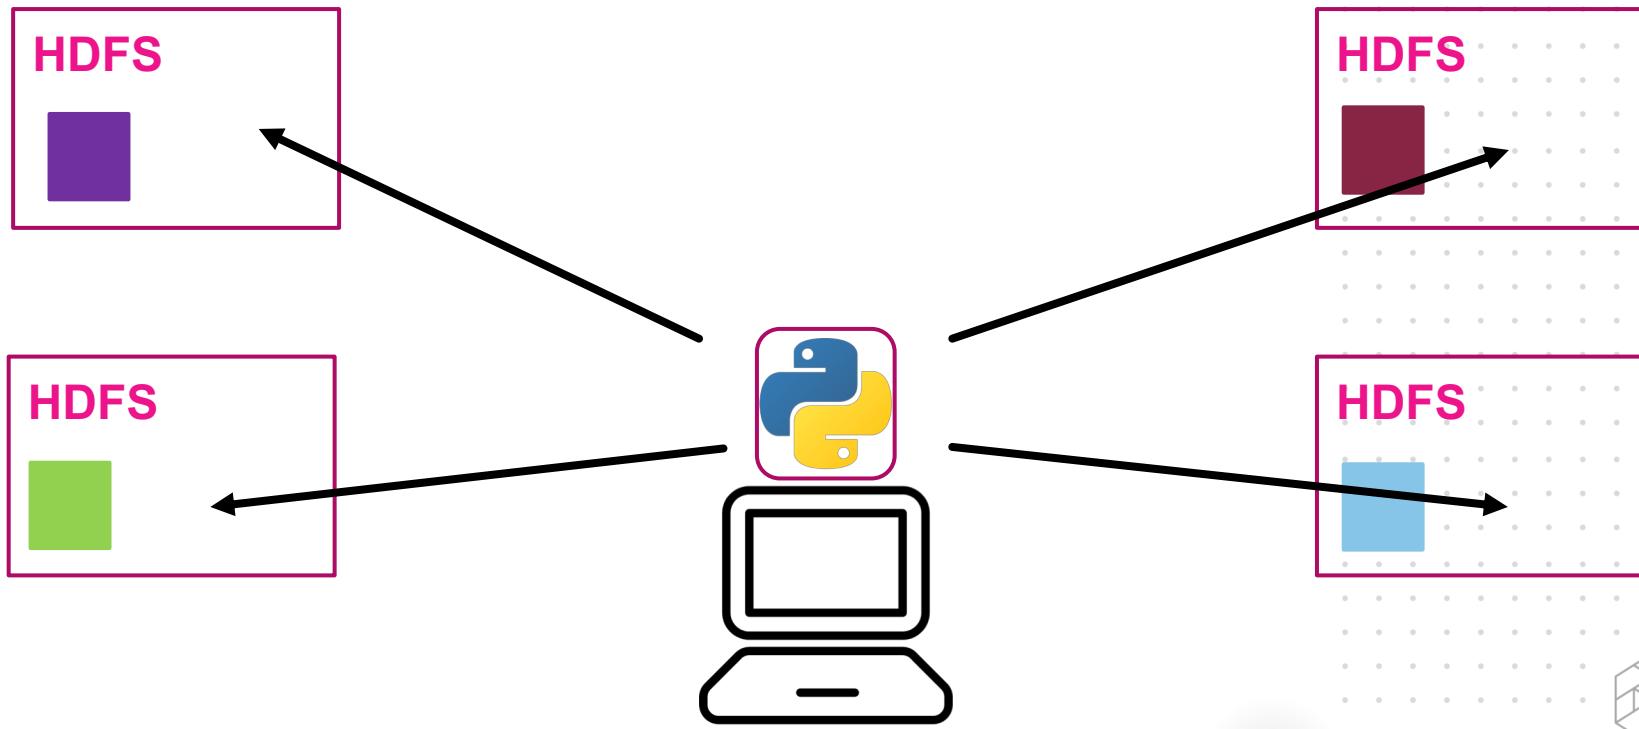
HDFS



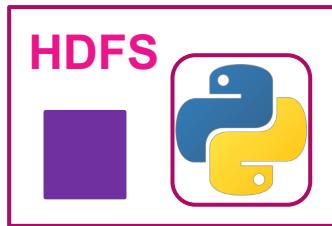
Pull-and-Process



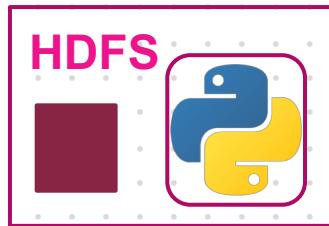
Use the Cluster



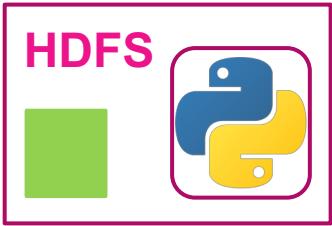
Use the Cluster



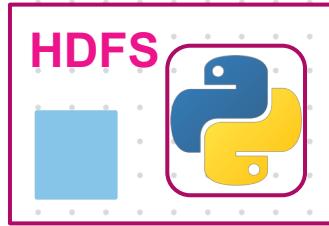
MY DATA SAYS 4.



MY DATA SAYS 3.



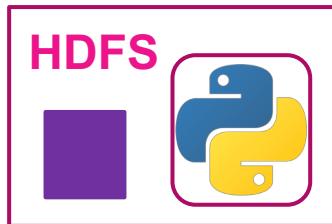
MY DATA SAYS 5.



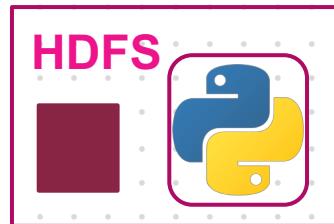
MY DATA SAYS 7.



Use the Cluster

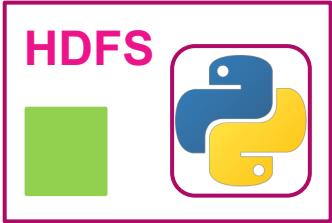


MY DATA SAYS 4.

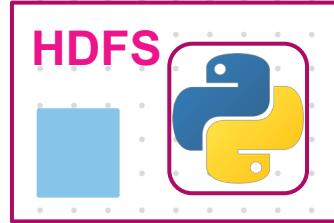


MY DATA SAYS 3.

This is called a
Mapping Stage



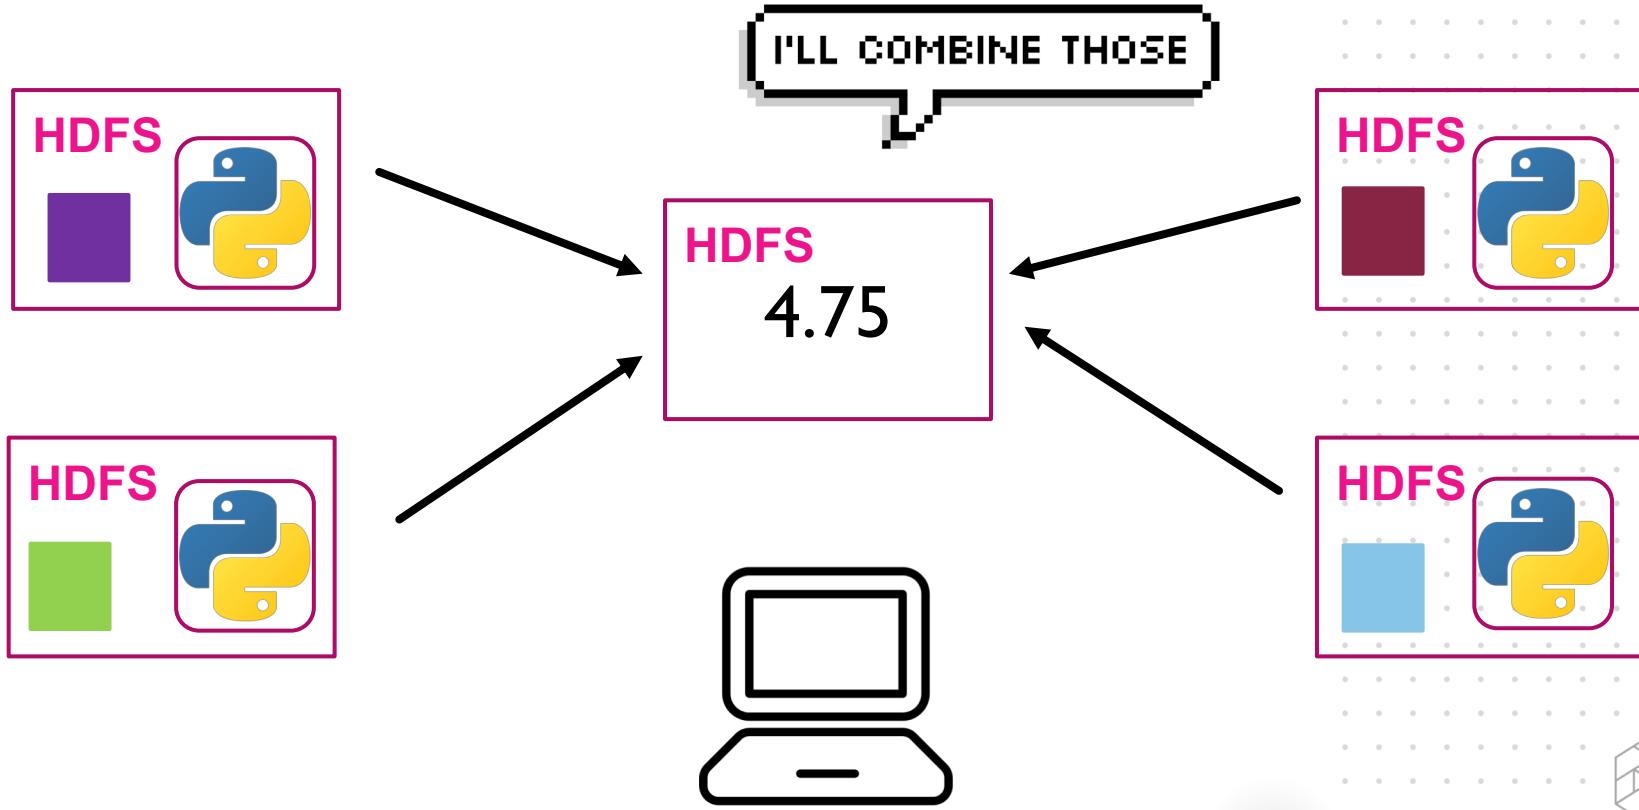
MY DATA SAYS 5.



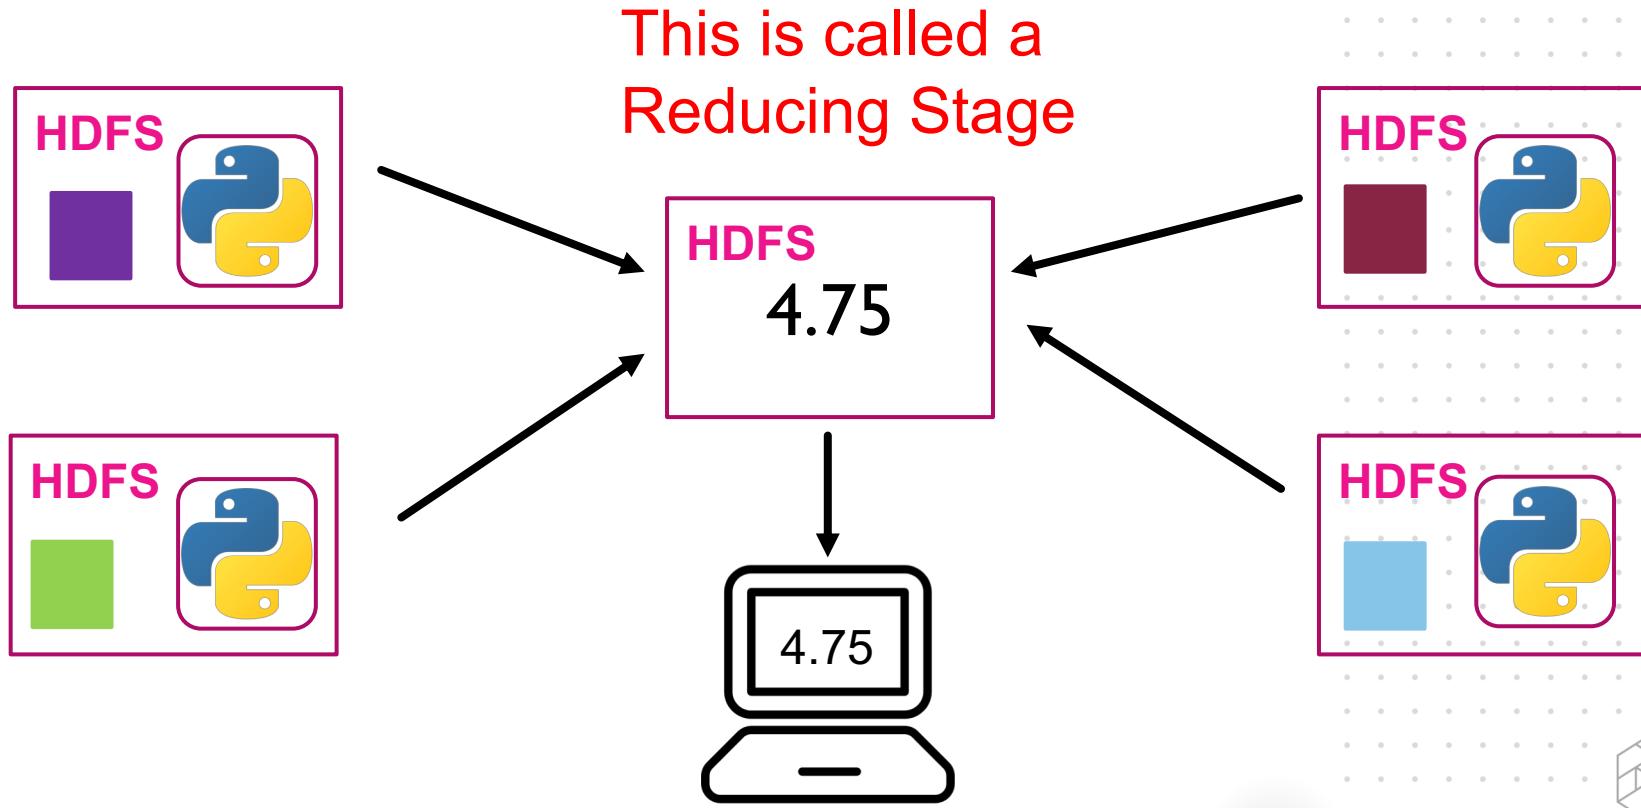
MY DATA SAYS 7.



Use the Cluster



Use the Cluster



MapReduce



- The “ah-ha” of MapReduce is to push the work to the data
- Transferring scripts is easy. Transferring data is hard.
- “Convert a big problem into lots of small problems”



Map Phase



- Processes data on each node and outputs to a collection stream.
- Can be anything that has consistent output
- Can work on rows or blocks or raw text, etc



Reduce Phase



- Takes all the outputs from the Mappers and combines them into a condensed output
- Could be averaging or counting or filtering or whatever you can code
- Passes output to either next reducing phase or back to the user

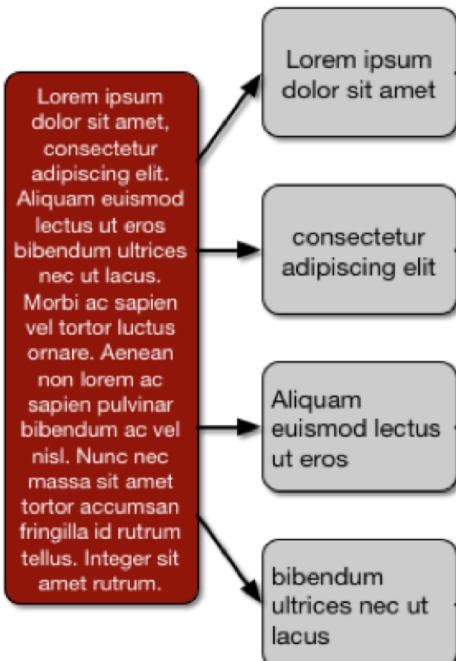


USING MAPREDUCE TO COUNT WORDS



Unstructured
Data Input

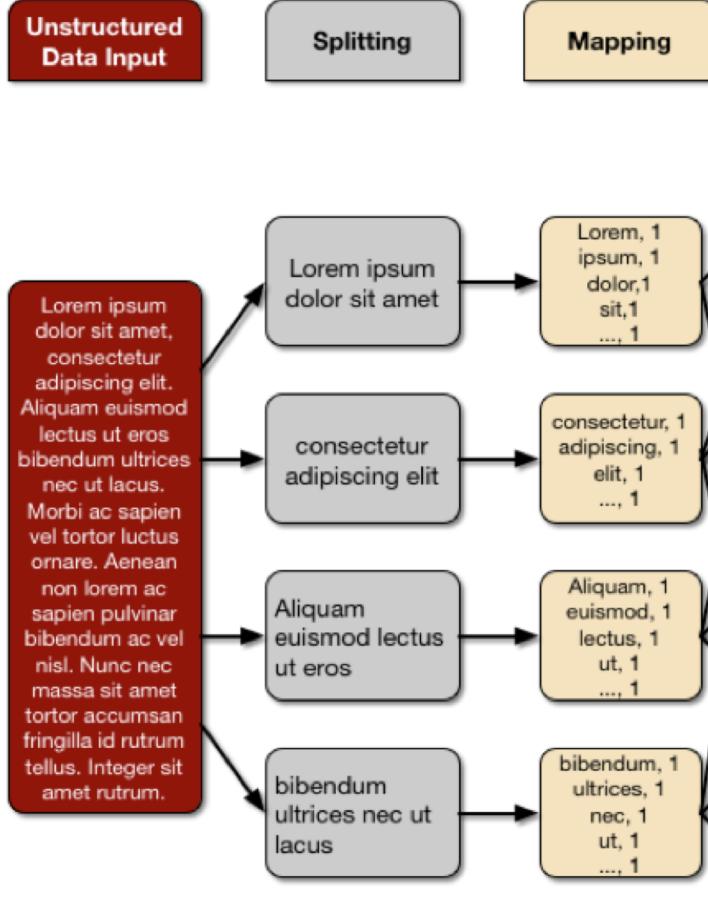
Splitting



We load our data onto HDFS and have it chunked up into different DataNodes.

If we want to count the total number of occurrences for each word, how could we do it with mapping?





We have each a program that writes each word into a stream. It's traditional to do it in key-value pairs. So we see:

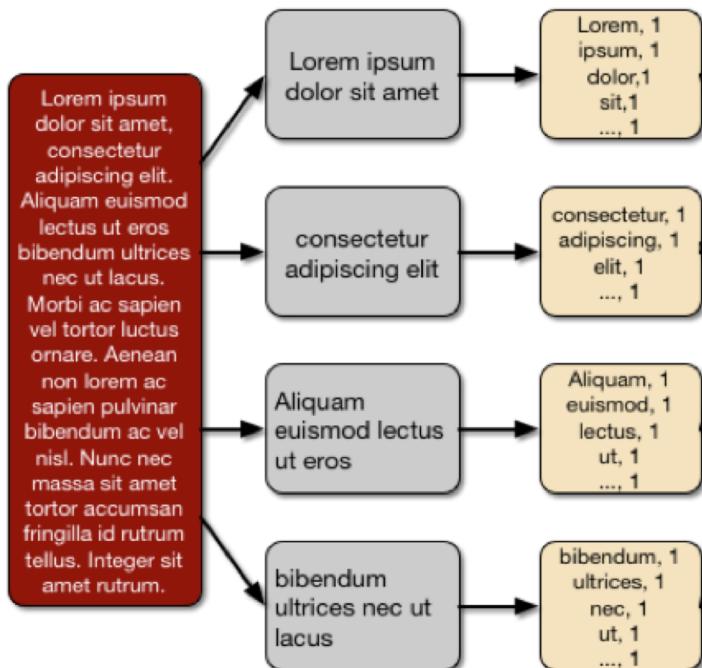
`word, 1 (appearance)`



Unstructured Data Input

Splitting

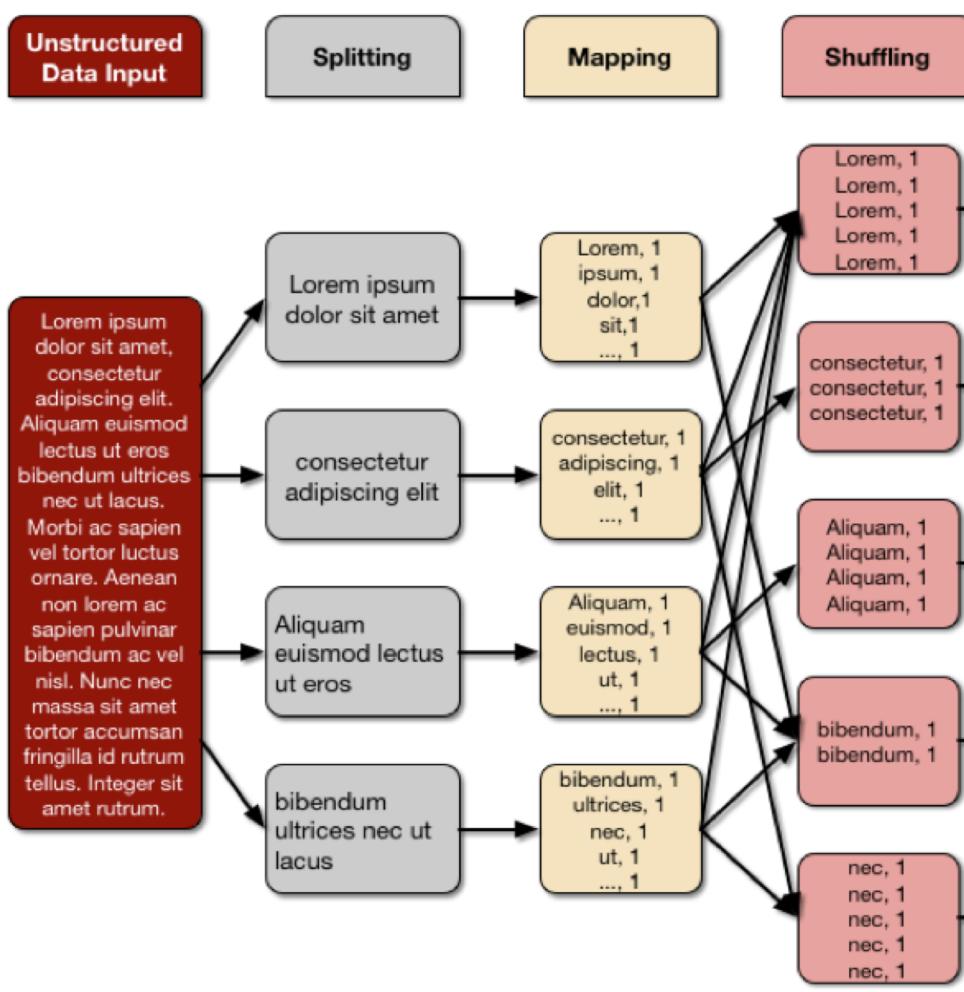
Mapping



One of the keys to MapReduce is to write data only to the stream. No objects that require long-term memory, so we can't use dictionaries to count.

So how can we reduce?

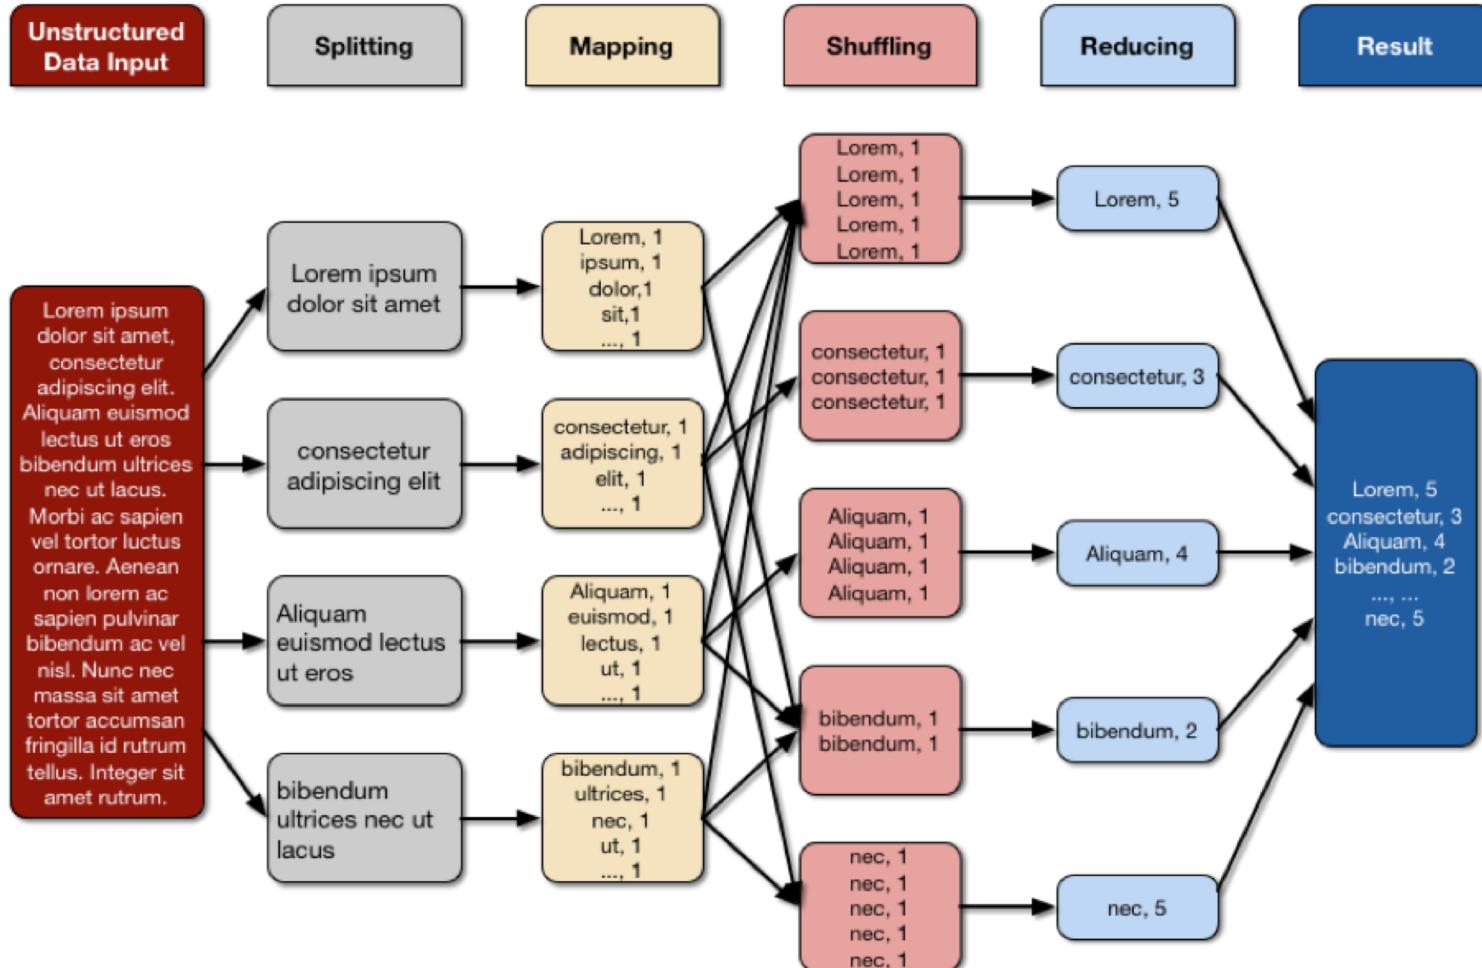




Sometimes we sort the data via a “shuffle” stage to help the reducer be more efficient.

Now that the data looks like this, how could we reduce?





When should I use Hadoop?



- Massive data. Must be at multiple terabytes or it simply isn't worth the overhead of maintaining a cluster.
- Hadoop isn't a “machine learning tool” it's about maintaining, pipelining, and cleaning data. We can make it do ML, but remember that it's designed for ETL.
- MapReduce reads from hard drive, so it can work, but it's slower than from RAM.



What else should I know?



- There's a whole suite of Hadoop based tools.
- Hive allows SQL queries to become MapReduce programs automatically
- Pig is another processing library that looks a lot like SQL
- Sqoop allows you to load data from SQL databases to HDFS
- Etc, etc, etc. This is a well supported tool set and if you need something, someone has likely built it.



HIVE



- Hive auto-converts SQL queries into MapReduce programs and can run on the distributed data
- Can work with both structured and unstructured data using it's special version of SQL called HiveQL
- Extremely robust to faults and great for exploring data
- Essentially allowed the use of Hadoop as an analytics tool



What else should I know?



- However, almost all of the Hadoop tools are now outperformed by Spark.
- So the main legacy of Hadoop is HDFS and its management tools like YARN. You don't need to understand those deeply – but know that HDFS is special because almost all the big data tools use it (like Spark).
- Understanding Hadoop and the Map-Reduce pattern forms the basis for all future big data technologies, so take this to heart.





QUESTIONS?



Cites



- Pixel speech bubbles: <https://pixelspeechbubble.com/>

